

What a leaderboard of clinical-note models can and cannot tell you

Sixteen models wrote psychotherapy session notes on two published protocols, and two independent LLM judges scored every one of them. This is what the exercise says about choosing a model — and about reading anybody's leaderboard, including this one.

therapy-note-bench · harness 0.2.0 · every figure here is generated from docs/leaderboard.json ·

<https://jannehyba.github.io/therapy-note-bench/>

100/101

pairs, agreed

Where both judges can tell two systems apart, they agree on which is better in 100 of 101 pairs. The judges are not the problem.

13/19

changed place

And this many systems still land somewhere else. A leaderboard position is not a measurement; it is a place in a queue, and a few real moves renumber everyone below them.

±0.02

each judge's own
vendor

Both judges score their own vendor's models higher by about this much, and the evidence cannot rule out zero for either. If you grade models with a model, measure this — and resample the models, not just the cases.

0.00–0.55

the production gap

Every model reliably records what happened last session. Almost none can say what happens at the next one. That is a feature gap, not a scoring artefact.

Read this if you are building or buying one of these

Two datasets exist for turning a therapy transcript into a clinical note. Both were benchmarked once, on models from 2024 and 2025, and never re-run. This benchmark re-runs them on 16 current models with the original prompts and the original scoring protocols, reproduced word for word so the numbers mean what the papers' numbers meant.

Nothing here is a clinical validation. No model in this benchmark has been shown to write a note a clinician would sign. What is measured is narrower and checkable: how much of a published rubric a note covers, how much of a published form it fills, and how far two different judges agree about either.

What was measured, and on what

Two tracks, because the two published protocols ask different questions and disagree about the answers — which is itself a result the source papers reported.

TRACK	WHAT THE MODEL WRITES	HOW IT IS SCORED	SESSIONS
TN-Eval SOAP	One SOAP note per conversation, from a transcript	A judge answers 23 yes/no rubric questions about it, plus one rating per sentence and one for factual accuracy. No gold note is involved.	50
iCARE / iHOPE	A 17-section clinical form, one call per section	Word overlap and embedding similarity against an expert-written note, plus five judge-rated dimensions and two time-bearing sections counted separately.	40

Both corpora are transcripts of published counselling demonstrations, not clinical sessions. On the iCARE side the experts themselves left most of the form empty — 316 of 680 fields say anything at all (46%). A model scores on those sections by staying quiet, so read the low rows as no signal rather than as a hard test.

The judge is a model, so the judge is measured first

TN-Eval released 150 notes that two trained therapists had already rated. Every candidate judge answers the same questions about the same notes before it is allowed near the leaderboard, and the agreement is published whatever it says.

MEASURE	FORM	JUDGE VS THERAPIST	THERAPIST VS THERAPIST
completeness	checklist	0.60	0.50
completeness	1–5 rating	0.11	0.13
conciseness	1–5 rating	0.03	0.19
faithfulness	1–5 rating	0.09	0.18

The right-hand column is the ceiling, not a target. **Two trained therapists rating the same notes barely agree on the 1–5 scales** — 0.13, 0.19 and 0.18, where 1.00 is perfect and 0 is chance. On the 23-item checklist they reach 0.50, and the judge reaches 0.60.

That is why the ranking uses the checklist and nothing else. The other three columns are reported because the protocol produces them, and ordering anything by them would be ordering by noise.

Measured across 150 of those notes, judge `gemini-3.1-pro-preview`. Krippendorff’s alpha, the statistic the source paper used, so the comparison is with their finding rather than with a re-definition of it.


One: a leaderboard position is not a measurement

The two judges see the same notes, ask the same 23 questions and reach almost the same conclusions about which of any two systems is better. And the two orderings they produce are visibly different. Both of those are true, and the reason is arithmetic rather than disagreement: nineteen systems are packed into a range of 0.22, so a handful of genuine moves rennumbers everyone below them.

What to do with that. Read a leaderboard as bands, not as an order. “In the top group” is a claim this kind of evidence supports; “ninth rather than tenth” is not, on this benchmark or on anybody else’s.

Two judges agree on 100 of the 101 pairs they can separate — and 13 of 19 systems still change place

Same notes, same rubric, same 23 criteria. Ordered by Completeness. A position is a place in a queue, not a measurement.



Grey lines held their place. A pair counts as separable when a paired bootstrap puts the two systems in different groups under that judge — over 25 shared conversations for judge A and 42 for judge B, because more of A’s notes went unfinished. Systems that print the same score share a place.

Source: docs/leaderboard.json and docs/saturation-*.json, harness 0.2.0.

Completeness under each of the two judges. Grey lines held their place. The top group holds 4 systems under one judge and 5 under the other; 4 are in both.

Two: the judge has a vendor, and it shows

Both judges here also *write* notes in this benchmark, so each is marking some of its own homework. The size of that is measurable: take the difference between the two judges’ scores for each system, and compare the judges’ own vendors against the systems neither of them built.

JUDGE	ITS VENDOR	EFFECT	95% INTERVAL	DETECTED
gemini-3.1-pro-preview	google	+0.018	-0.011 to +0.048	no
gpt-5.6-terra	openai	+0.027	-0.002 to +0.059	no

Neither is detected once the models are treated as a sample rather than as the whole vendor. The units are completeness, so an effect of about 0.02 is a fiftieth of the range the models occupy — enough to move a system several places in an ordering this tight, and not enough to make a bad note look good. Both point estimates are positive and neither interval clears zero.

What to do with that. If you evaluate models with a model, run this check. Two judges from two vendors is the cheapest way to have it; one judge cannot measure its own bias at all, and a caveat in the methods section is not a measurement.


Two systems carry most of it, and they are the two that a definition change moved into these groups on 2026-08-26: dropping gemma4 takes the Gemini figure from +0.018 to +0.008, and dropping gpt-oss-120b takes the GPT one from +0.027 to +0.018. Nothing here tests whether the two judges differ from each other; asked directly, they do not — +0.010 [−0.015, +0.033].

Three: nothing here can tell you what happens next

The iCARE form has two time-bearing sections: what happened at the previous session, and what happens at the next one. Every model fills the first. Almost none fills the second, and the best of them manages it in just over half the sessions where an expert did.

Every model can say what happened last time. Almost none can say what happens next.

iCARE's two time-bearing sections, scored on the same 40 sessions. The fraction of the expert-answered sections each model also answered.



The forward column rests on eleven sessions, so a score of 0.09 there is one session and the gap between two models a few hundredths apart is not evidence. The backward column rests on thirty-four. They were one averaged column until harness 0.2.0, and the average — weighted three to one towards the easy half — published the opposite of the finding this track exists to reproduce.

Source: docs/leaderboard.json, iCARE track, judge gemini-3.1-pro-preview, harness 0.2.0. Neither column involves the judge: both are computed from the note and the expert note alone.

The fraction of expert-answered sections each model also answered. Backward on the left of each pair, forward on the right.

What to do with that. If your product promises continuity between sessions — a plan, a follow-up, something to pick up next time — that is the part no current model produces on its own, and it is the part a demo will not show you because a demo is one session.

How much room is left

A benchmark that everything passes has stopped measuring. This one has not, and it has not evenly: under gemini-3.1-pro-preview the twenty-three rubric criteria are in four different states at once.

CRITERION	HOW MANY	WHAT THAT MEANS
Every model already does it	3 of 23	Free points. They raise every score by the same amount and separate nobody.
Nobody does it, the therapist included	2 of 23	Dead. These transcripts do not contain the answer, so the criterion measures the corpus rather than the model.
Still separates models	13 of 23	Where the benchmark is doing its job.
Partly	5 of 23	Separates some models and not others.

These are one judge’s verdicts. A verdict reads that judge’s answers, and the second judge does not return the same reading: gpt-5.6-terra puts 3 criteria on the floor rather than 2, the extra one being assessment tools, which the first judge leaves 0.02 above it. docs/limitations.md excludes only the 2 both of them put there.

Strip out the 2 nobody reaches and the most a model could score is 0.93 — weighted the way completeness is, four sections averaged equally rather than 21 criteria out of 23. **The best model here reaches 0.546, which is 59% of that.** There is room.

How fast it is being used up: the two 2025 models the source paper benchmarked score 0.364 and 0.389; the 16 current ones span 0.446 to 0.546. **One model generation moved the top of the table by +0.157.** Two or three more at that rate and the reachable part of this rubric is exhausted — which is a reason to record what the corpus and the protocol are now, not a reason to trust the ranking more.

The other track’s judge-scored measure is nearly out of room

TRACE rates five dimensions from 1 to 5. Under one judge all 16 models land between 4.76 and 4.99 — 6% of the scale. Under the other, 3.58 to 4.11, which is 13%. The two judges’ orderings correlate at +0.83 and place 11 of 16 systems differently anyway.


What to do with that. A measure where every model scores within a few percent of every other is not evidence that the models are equally good; it is evidence that the measure is running out. That the two judges disagree about how much room is left — 6% against 13% — is a fact about the judges, and it is why one judge is not enough to notice this happening. If you are building an evaluation, the per-criterion breakdown is the thing to watch — an aggregate stays healthy for a long time after the parts of it have died.

Computed by a paired bootstrap over the conversations every system was scored on, published as docs/saturation-<judge>.json and drawn in full on the methods page. TRACE has no human anchor at all: unlike the rubric, no therapist ever rated these notes on it, and it is labelled that way wherever it appears.

What these numbers are not

Every model covers more of the rubric than the therapist does — which is not the same as writing a better note

Completeness: the fraction of 23 rubric criteria a judge found present. Judge gemini-3.1-pro-preview.




A therapist writes what matters for the next session and leaves out what does not; the rubric counts what is present and cannot see why anything was left out. TN-Eval reported the same direction from blinded expert comparison, so this is a reproduction rather than an anomaly — and it is a statement about a checklist, not about clinical writing. Quote the number with this sentence attached, or do not quote it. Source: docs/leaderboard.json, harness 0.2.0. The same ordering holds under the second judge; the figures differ.

Completeness, every system, under the first judge. The therapist-written note is the bottom row.

Every model covers more of the rubric than the therapist does. That is a reproduction — the source paper found the same direction from blinded expert comparison — and it is a statement about a checklist. A therapist writes what matters for the next session and leaves out what does not; the rubric counts what is present and cannot see why anything was left out.

The one exception is the measure with the weakest human agreement: on factual accuracy, one of the two 2025 reference models scores *below* the therapist under both panel judges and above her under a third. A column that changes sign when the referee changes is not measuring what the other two are.

Each dot is one system. Horizontal: Completeness. Vertical: Faithfulness. The dashed line is least squares, drawn rather than described.



Source: docs/leaderboard.json, harness 0.2.0. Spearman as computed for the concordance panel.

Neither track measures whether a note is clinically useful, safe to put in a record, or acceptable to a supervisor. Nobody has run that study on these models. Treat every figure here as a lower bound on what you would have to check yourself.

How to check any of this

Every figure and every number in this document is generated from two published files. Nothing was typed in, and nothing here can drift from the site without the tests failing.

And four pages carry what the numbers rest on: [the datasets](#) — where each came from, what licence it publishes (two of the three publish none) and the traps in them; [the method](#); [what a result cannot claim](#); and [what exists in this field](#) and what does not.

- `docs/leaderboard.json` — every row, every measure, and the six version fields a row has to agree on before it may be compared with another.
- `docs/saturation-<judge>.json` — the paired bootstrap that decides which systems this evidence can tell apart.
- `results/rows.jsonl` — append-only. A re-run adds rows beside the old ones; what is drawn is the newest of each.

The figures redraw with `make figures` and this document with `make brief`. There are 5 of them, and each carries the file it came from in its own footer.

What a result here cannot claim

- Two judges are two instruments. A number scored by one is never averaged with a number scored by the other, and a table that cannot say what settings its judge ran at is withdrawn rather than drawn.
- An absence is never counted as a zero. A note the judge did not finish is left out of the average and said so, rather than dragging a system down for a question nobody answered.
- The corpora are demonstrations, not clinical sessions, and two of the three sources publish no licence at all. Nothing is redistributed here.