What a leaderboard of clinical-note models can and cannot tell you

Sixteen models wrote psychotherapy session notes on two published protocols, and two independent LLM judges scored every one of them. This is what the exercise says about choosing a model — and about reading anybody's leaderboard, including this one.

therapy-note-bench · harness 0.6.0 · scored 2026-08-27 · every table and figure here is generated from the files the site publishes; the figures in the prose are listed in `tests/test_brief.py` · https://jannehyba.github.io/therapy-note-bench/

100/101
pairs, agreed

Where both judges can tell two systems apart, they agree on which is more complete in 100 of 101 pairs — one measure on one track. The judges are not the problem.

13/19
changed place

And this many systems still land somewhere else. A leaderboard position is not a measurement; it is a place in a queue, and a few real moves renumber everyone below them.

+0.02
each judge's own vendor

Both judges score their own vendor’s models higher by about this much, and the evidence cannot rule out zero for either. If you grade models with a model, measure this — and resample the models, not just the cases.

0.00–0.55
the production gap

Every model writes something where the expert wrote about the last session. Almost none writes anything about the next one, on the eleven sessions where an expert did. Both columns check that a field was answered, not that the answer is right.

Read this if you are building or buying one of these

Two datasets exist for turning a therapy transcript into a clinical note. Both were benchmarked once, on models of the 2023–2024 era, and never re-run. This benchmark re-runs them on 16 current models with the generation prompts reproduced word for word. On the SOAP track the scoring protocol is the paper’s too. On the iCARE track it is not, and this page says so where those numbers appear: the word-overlap column is computed on the fields the expert answered and is not comparable with the paper’s table, and TRACE is a re-implementation with a prompt of our own.

Both corpora are English, and nothing here says anything about a session held in another language. The 16 models are what one endpoint and the two judges’ own vendors had deployed on the run date — not a survey of the market, and nothing outside that deployment was tested.

Nothing here is a clinical validation. No model in this benchmark has been shown to write a note a clinician would sign. What is measured is narrower and checkable: how much of a published rubric a note covers, how much of a published form it fills, and how far two different judges agree about either.

What was measured, and on what

Two corpora and four published instruments, because the corpora and the protocols on them ask different questions and disagree about the answers — which is itself a result the source papers reported.

TRACK	WHAT THE MODEL WRITES	HOW IT IS SCORED	SESSIONS
TN-Eval SOAP	One SOAP note per conversation, from a transcript	A judge answers 23 yes/no rubric questions about it, plus one rating per sentence and one for factual accuracy. No gold note is involved. The same notes are also rated on eight of PDSQI-9's nine attributes — an instrument validated on discharge summaries rather than on session notes, with no human anchor here, and eight of the eleven columns the published table draws. The two are never averaged: they are separate instruments, and the leaderboard is ordered by the rubric alone.	50
iCARE / iHOPE	A 17-section clinical form, one call per section	Word overlap and embedding similarity against an expert-written note, plus TRACE — five dimensions rated 1 to 5, re-implemented here because the authors published the dimension names and never a prompt — and two time-bearing sections counted separately.	40

Both corpora are transcripts of published demonstration sessions, not of clinical practice. On the iCARE side the experts themselves left most of the form empty — 316 of 680 fields say anything at all (46%). On the overlap and time-bearing columns those blank fields are out of the denominator, so nothing is scored on them either way; TRACE rates the whole rendered note, so padding an empty field is still read there. Read a high iCARE score as "fills the form correctly, including knowing when not to", not as "writes a good clinical note".

The judge is a model, so the judge is measured first

TN-Eval released 150 notes that two trained therapists had already rated. Every candidate judge answers the same questions about the same notes before it is allowed near the leaderboard, and the agreement is published whatever it says.

MEASURE	FORM	STATISTIC	JUDGE VS THERAPIST	THERAPIST VS THERAPIST	ALPHA, JUDGE	ALPHA, THERAPISTS
completeness	checklist	Cohen's kappa	0.60	0.50	0.60	0.50
completeness	1–5 rating	Spearman rho	0.24	0.13	0.11	0.13
conciseness	1–5 rating	Spearman rho	0.11	0.19	0.03	0.19
faithfulness	1–5 rating	Spearman rho	0.09	0.18	0.09	0.18

Both therapist-vs-therapist columns are a ceiling, not a target. **Two trained therapists rating the same notes barely agree on the 1–5 scales** — 0.13, 0.19 and 0.18 by alpha, where 1.00 is perfect and 0 is chance. On the 23-item checklist they reach 0.50, and the judge reaches 0.60.

Why two pairs. Cohen's kappa suits a yes/no criterion and Spearman's rho an ordered scale, so those are the statistics the leaderboard and the methods page report. Krippendorff's alpha sits beside them because it is what TN-Eval used, and a comparison with their finding has to be on their statistic. The two do not agree: on likert completeness the judge scores 0.24 by rho against a ceiling of 0.13, and 0.11 by alpha against a ceiling of 0.13 — above the therapists on one and below them on the other. Which statistic is quoted decides the answer, so both are printed.

That is why the ranking uses the checklist and nothing else. The other three columns are reported because the protocol produces them, and ordering anything by them would be ordering by noise.

Measured across 150 of those notes, judge `gemini-3.1-pro-preview`. The second judge, `gpt-5.6-terra`, clears the same ceiling by less: 0.52 against 0.50 on the checklist. Both are in `docs/judges.json`.

One: a leaderboard position is not a measurement

The two judges see the same notes, ask the same 23 questions and reach almost the same conclusions about which of any two systems is better. And the two orderings they produce are visibly different. Both of those are true, and the reason is arithmetic rather than disagreement: 19 systems are packed into a range of 0.22, so a handful of genuine moves renumbers everyone below them.

What to do with that. Read a leaderboard as bands, not as an order. “In the top group” is a claim this kind of evidence supports; “ninth rather than tenth” is not, on this benchmark or on anybody else’s.

Two judges agree on 100 of the 101 pairs they can separate — and 13 of 19 systems still change place

Same notes, same rubric, same 23 criteria. Ordered by Completeness. A position is a place in a queue, not a measurement.

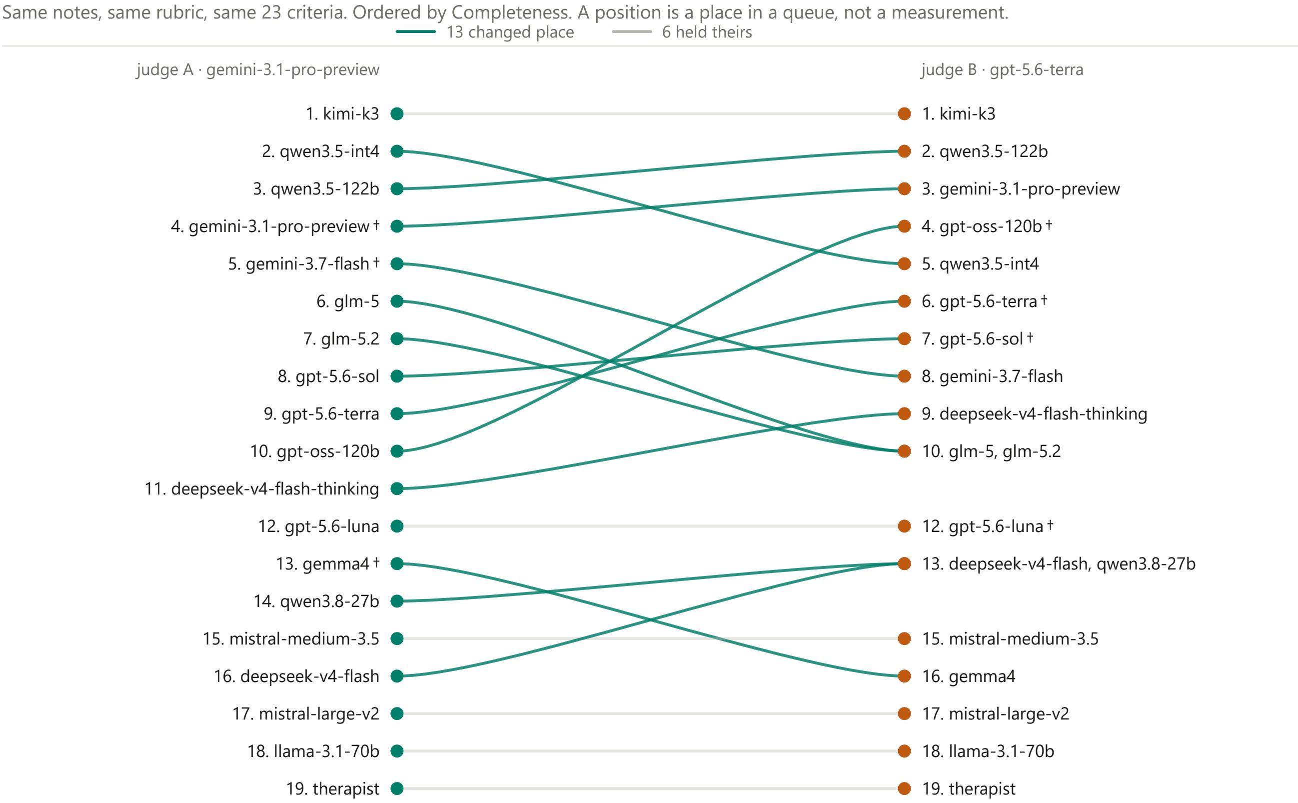

Grey lines held their place. A pair counts as separable when a paired bootstrap puts the two systems in different groups under that judge — over 25 shared conversations for judge A and 42 for judge B, because more of A's notes went unfinished. Systems that print the same score share a place, and a shared place is drawn as one label. A dagger marks a system the judge on that side shares a vendor with.
Source: docs/leaderboard.json and docs/saturation-*.json, harness 0.6.0.

Completeness under each of the two judges. Grey lines held their place. The top group holds 4 systems under one judge and 5 under the other; 4 are in both.

Two: the judge has a vendor, and nobody can rule out that it shows

Both judges here also *write* notes in this benchmark, so each is marking some of its own homework. The size of that is measurable: take the difference between the two judges' scores for each system, and compare the judges' own vendors against the systems neither of them built.

JUDGE	ITS VENDOR	EFFECT	95% INTERVAL	DETECTED
`gemini-3.1-pro-preview`	google	+0.018	−0.010 to +0.047	no
`gpt-5.6-terra`	openai	+0.027	−0.004 to +0.058	no

Neither is detected once the models are treated as a sample rather than as the whole vendor. The units are completeness, so these effects are 18% and 26% of the range the current models occupy under their own judge — enough to move a system several places in an ordering this tight, and not enough to make a bad note look good. Both point estimates are positive and neither interval clears zero.

What is behind them: 3 of google's own models against 12 neither judge wrote, over 25 conversations; 4 of openai's own models against 12 neither judge wrote, over 25 conversations. Not detected is not the same as absent, and a comparison this small is one reason why.

What to do with that. If you evaluate models with a model, run this check. Two judges from two vendors is the cheapest way to have it; one judge cannot measure its own bias at all, and a caveat in the methods section is not a measurement.

Each effect leans on one system more than the others, and they are the two that a definition change moved into these groups on 2026-08-26: dropping `gemma4` takes the Gemini figure from +0.018 to +0.008, and dropping `gpt-oss-120b` (scored on 42 of 50 conversations; the rest were lost to an output format the rubric parser truncates) takes the GPT one from +0.027 to +0.018. Nothing here tests whether the two judges differ from each other; asked directly, they do not — +0.009 [−0.039, +0.058].

Three: nothing here can tell you what happens next

The iCARE form has two time-bearing sections: what happened at the previous session, and what happens at the next one. Every model reliably fills the first — fourteen of sixteen on all thirty-four sessions the experts answered, the other two on thirty-three. Almost none fills the second, and the best of them manages it in just over half the sessions where an expert did.

Every model can say what happened last time. Almost none can say what happens next.

iCARE's two time-bearing sections, scored on the same 40 sessions. The fraction of the expert-answered sections each model also answered.

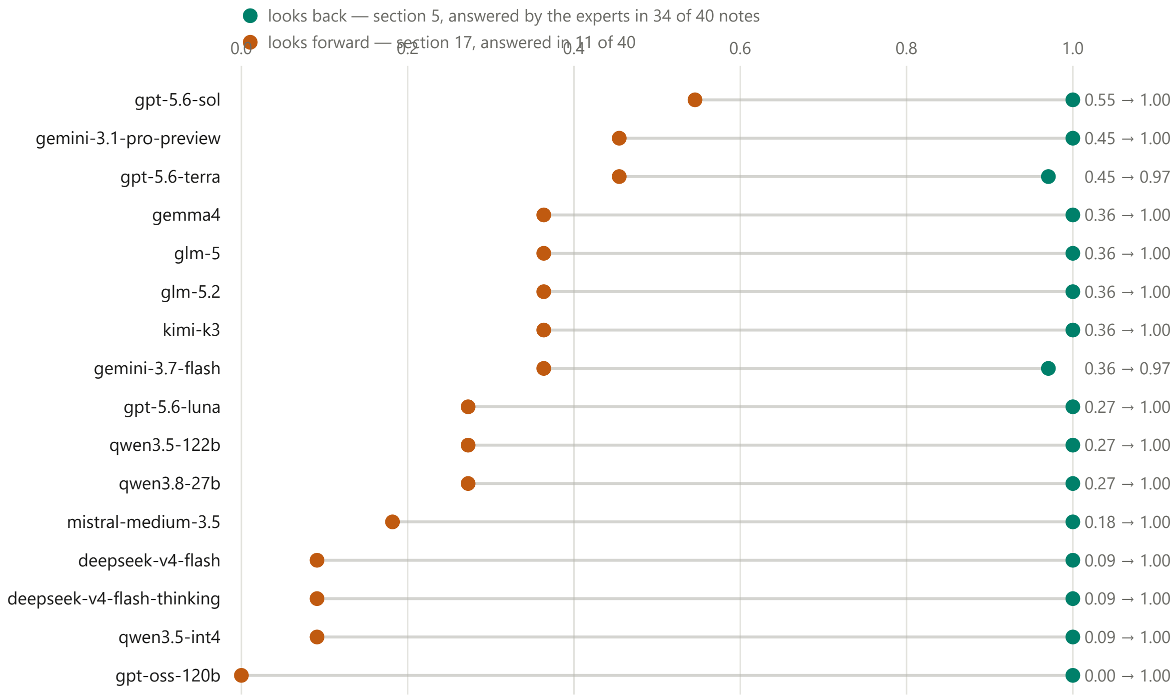

The forward column rests on eleven sessions, so a score of 0.09 there is one session and the gap between two models a few hundredths apart is not evidence. The backward column rests on thirty-four. They were one averaged column until harness 0.2.0, and the average — weighted three to one towards the easy half — published the opposite of the finding this track exists to reproduce.

Source: docs/leaderboard.json, iCARE track, judge gemini-3.1-pro-preview, harness 0.6.0. Neither column involves the judge: both are computed from the note and the expert note alone.

The fraction of expert-answered sections each model also answered. Looking back and looking forward, one pair per model; which dot is which is in the legend, not in their order.

What to do with that. If your product promises continuity between sessions — a plan, a follow-up, something to pick up next time — that is the part no current model produces on its own, and it is the part a demo will not show you because a demo is one session.

How much room is left

A benchmark that everything passes has stopped measuring. This one has not, and it has not evenly: under `gemini-3.1-pro-preview` the twenty-three rubric criteria are in four different states at once.

CRITERION	HOW MANY	WHAT THAT MEANS
Every model already does it	3 of 23	Free points. They raise every score by the same amount and separate nobody.
Nobody does it, the therapist included	2 of 23	Dead. These transcripts do not contain the answer, so the criterion measures the corpus rather than the model.
Still separates models	13 of 23	Where the benchmark is doing its job.
Partly	5 of 23	Separates some models and not others.

These are one judge's verdicts. A verdict reads that judge's answers, and the second judge does not return the same reading: `gpt-5.6-terra` puts 3 criteria on the floor rather than 2, the extra one being assessment tools, which the first judge leaves 0.02 above it. `docs/limitations.md` excludes only the 2 both of them put there.

Strip out the 2 nobody reaches and the most a model could score is 0.93 — weighted the way completeness is, four sections averaged equally rather than 21 criteria out of 23. **The highest completeness here is 0.546, which is 59% of that.** There is room.

How fast it is being used up: the two 2024 models whose notes the source paper released score 0.364 and 0.389; the 16 current ones span 0.446 to 0.546. **One model generation moved the top of the table by +0.157.** Two or three more at that rate and the reachable part of this rubric is exhausted — which is a reason to record what the corpus and the protocol are now, not a reason to trust the ranking more.

The other track's judge-scored measure is nearly out of room

TRACE is a five-dimension framework from the iCARE paper, rated 1 to 5. Its authors published the dimension names and never a prompt, so this column is a re-implementation with a prompt of our own. Under one judge all 16 models land between 4.76 and 4.99 — 6% of the scale. Under the other, 3.58 to 4.11, which is 13%. The two judges' orderings correlate at +0.83 and place 11 of 16 systems differently anyway.

What to do with that. A measure where every model scores within a few percent of every other is not evidence that the models are equally good; it is evidence that the measure is running out. That the two judges disagree about how much room is left — 6% against 13% — is a fact about the judges, and it is why one judge is not enough to notice this happening. If you are building an evaluation, the per-criterion breakdown is the thing to watch — an aggregate stays healthy for a long time after the parts of it have died.

Computed by a paired bootstrap over the conversations every system was scored on, published as `docs/saturation-<judge>.json` and drawn in full on the methods page. TRACE is a re-implementation and has no human anchor at all: unlike the rubric, no therapist ever rated these notes on it, and it is labelled that way wherever it appears.

What these numbers are not

Every model covers more of the rubric than the therapist does — which is not the same as writing a better note

Completeness: the equal-weighted mean of a note's four section fractions over TN-Eval's 23 criteria, not the fraction of all 23. Judge gemini-3.1-pro-preview.

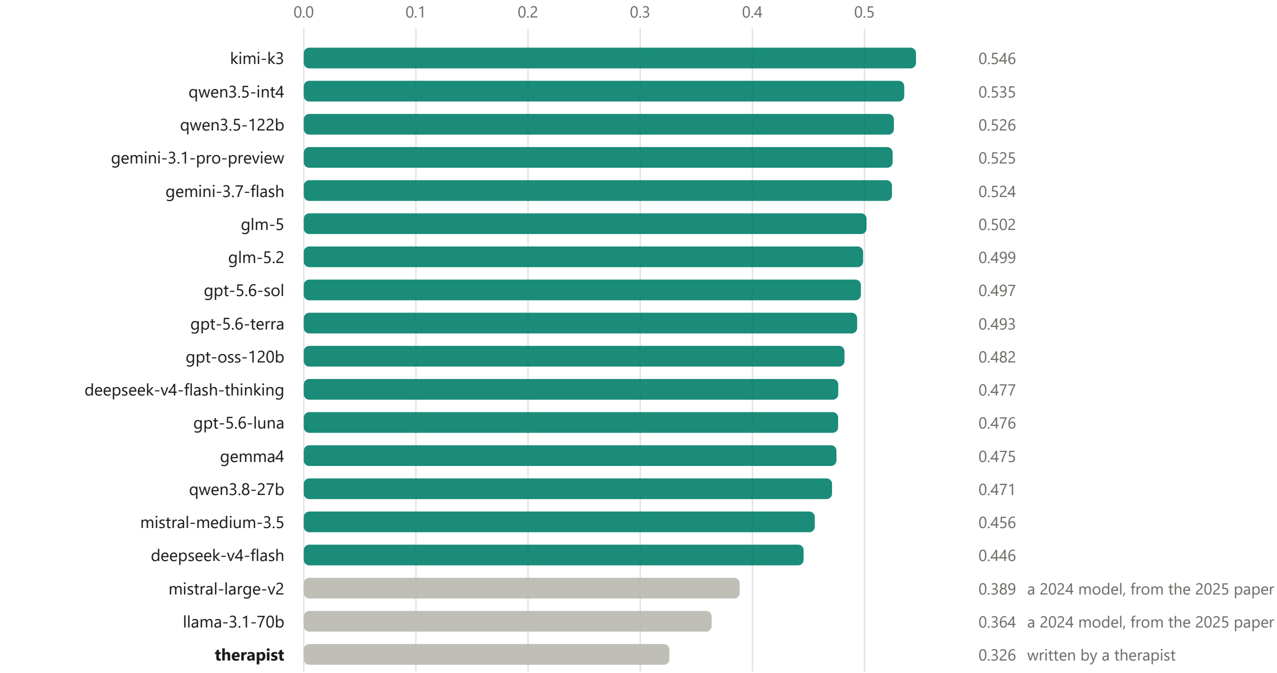

A therapist writes what matters for the next session and leaves out what does not; the rubric counts what is present and cannot see why anything was left out. TN-Eval reported the same direction from blinded expert comparison, so this is a reproduction rather than an anomaly — and it is a statement about a checklist, not about clinical writing. Quote the number with this sentence attached, or do not quote it.

Source: docs/leaderboard.json, harness 0.6.0. The same ordering holds under the second judge; the figures differ.

Completeness, every system, under the first judge. The therapist-written note is the bottom row.

Every model covers more of the rubric than the therapist does. That is a reproduction — the source paper found the same direction from blinded expert comparison — and it is a statement about a checklist. A therapist writes what matters for the next session and leaves out what does not; the rubric counts what is present and cannot see why anything was left out.

The one exception is the measure with the weakest human agreement: on factual accuracy, one of the two 2024 reference models scores *below* the therapist under both panel judges and above her under `gemini-2.5-pro` — a candidate judge measured at harness 0.2.0 and withdrawn from the tables, so that last comparison is not in the files this page is drawn from. A column that changes sign when the referee changes is not measuring what the other two are.

Answering more of the checklist and inventing more are not the same trade under the two judges

Each dot is one system. Horizontal: Completeness. Vertical: Faithfulness. The dashed line is least squares, drawn rather than described.

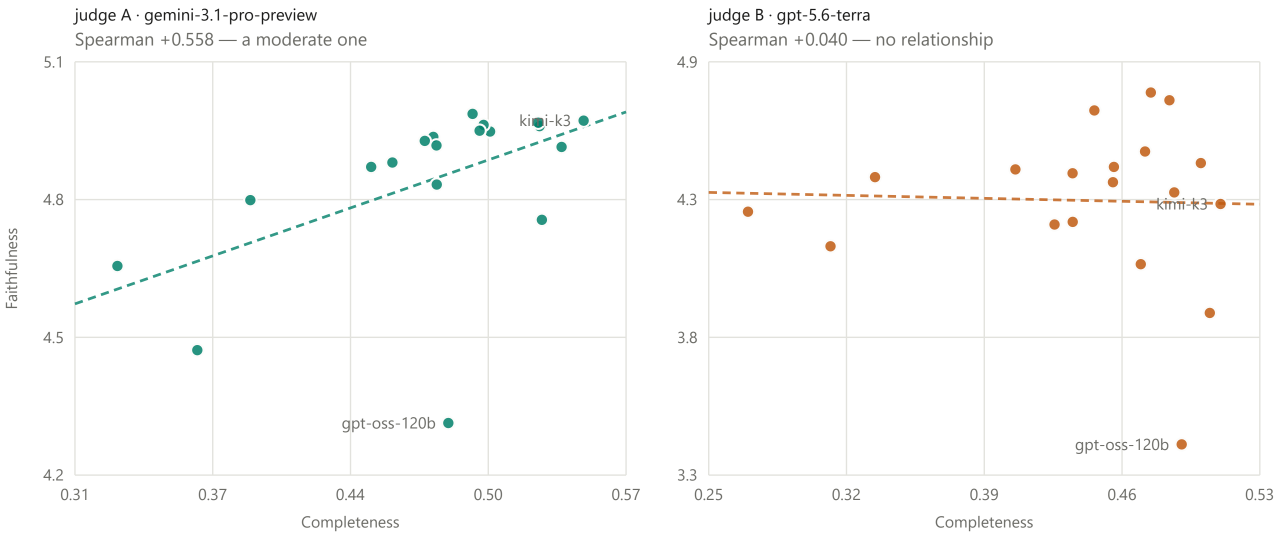

Faithfulness is the weakest column here: two trained therapists rating the same notes reach a Krippendorff's alpha of 0.18 on it. A relationship that appears under one judge and not the other is a fact about the instrument before it is a fact about the models.

Source: docs/leaderboard.json, harness 0.6.0. Spearman as computed for the concordance panel.

Completeness against factual accuracy, one panel per judge. The two judges do not agree about whether covering more of a checklist goes with inventing more. Only the two ends are named — nineteen labels in a panel this size collide, and the figure is about the contrast between the judges rather than about any one model. Every system is named in the chart above.

Neither track measures whether a note is clinically useful, safe to put in a record, or acceptable to a supervisor. Nobody has run that study on these models. Treat every figure here as a lower bound on what you would have to check yourself.

How to check any of this

Every figure in this document is drawn from the files the site publishes, and every table is built from them row by row. Five files rather than two: `leaderboard.json`, the two `saturation-*.json`, `calibration.json` and `corpus-profile.json`.

The prose around them is written by hand. Some figures in it are computed from the payload and a test fails if they drift; others are typed here beside the prose they belong to. The test names the sentences it covers, and a sentence it does not name is a sentence nobody is checking. This paragraph used to say that nothing here was typed in. That was false in this file more than thirty times, and it read as an instruction not to look.

And four pages carry what the numbers rest on: the datasets — where each came from, what licence it publishes (three of the six publish none for their data) and the traps in them; the method; what a result cannot claim; and what exists in this field and what does not.

- `docs/leaderboard.json` — every row, every measure, and the six version fields a row has to agree on before it may be compared with another.
- `docs/saturation-<judge>.json` — the paired bootstrap that decides which systems this evidence can tell apart.
- `results/rows.jsonl` — append-only. A re-run adds rows beside the old ones; what is drawn is the newest of each.

The figures redraw with `make figures` and this document with `make brief`. 5 are drawn and 4 of them are here — the fifth sits beside the saturation panel on the methods page — and each carries the file it came from in its own footer.

What a result here cannot claim

- Two judges are two instruments. A number scored by one is never averaged with a number scored by the other, and a table that cannot say what settings its judge ran at is withdrawn rather than drawn. **The rule is not bookkeeping.** Raising one judge's thinking budget from 128 to 256 tokens, with nothing else changed, moved every one of nineteen systems on completeness (mean +0.017) and conciseness (+0.048), changed six positions on the first and sixteen on the second, and reordered the conciseness top three. A setting no reader ever sees moved the table further than most of the differences printed in it. Those rows are history: they are not in this repository and cannot be re-derived from it, which is the other half of the lesson.
- An absence is never counted as a zero. A note the judge did not finish is left out of the average and said so, rather than dragging a system down for a question nobody answered.
- The corpora are demonstrations, not clinical practice. Of the six sources this benchmark draws on, two carry a licence; three publish none at all for their data, and the sixth shows an MIT badge on a code repository with no LICENSE file behind it. **No corpus is redistributed here** — the pages show scores, field names, TN-Eval's Apache-2.0 prompt and rubric, and one worked example quoting a single iHOPE note section in the clinician's wording and one model's.